

# CS747 - Assignment 1 Report

Shashank Kumar

September 25, 2020

---

## Directory Structure

- *bandit.py* - Main program file to run a bandit instance
- *run.py* - Wrapper to run the file *bandit.py* to generate outputs
- *plot.py* - Script to plot the generated outputs from *run.py*
- *epsilonGreedy.py* - Code to implement epsilon-greedy algorithm
- *ucb.py* - Implementation of UCB algorithm
- *ucb\_kl.py* - KL-UCB algorithm implementation
- *thompson\_sampling.py* - Code for regular Thompson sampling
- *thompson\_sampling\_hint.py* - Thompson sampling with the true means known
- *helper.py* - A file for additional tasks like reading and parsing instances

---

## Implementation Details

### Epsilon-Greedy

A number is drawn from a uniform distribution in the range  $[0, 1]$ , say  $\mathbf{n}$ . Based on the epsilon ( $\epsilon$ ) value specified, the arm with maximum empirical mean is chosen if the  $\mathbf{n} \leq (1 - \epsilon)$ . Otherwise an arm is chosen randomly from the bandit instance.

### UCB

The UCB means for all the arms are initialized to be  $\infty$ . For each run, the arm with maximum UCB mean is chosen and its mean gets updated based on its empirical mean, the time  $\mathbf{t}$ , and number of times it has been sampled  $u_a^t$ .

## KL-UCB

The KL-UCB equation is as follows:

$$ucb-kl_a^t = \max\{q \in [\hat{p}_a^t, 1] \text{ s.t. } u_a^t * KL(\hat{p}_a^t, q) \leq \ln(t) + c * \ln(\ln(t)) \text{ where } c \geq 3\}$$

For KL-UCB algorithm implementation, the constant  $c$  is chosen to be 3.

The amount by which the value of  $q$  is decreased after every step if the equation is not satisfied, is taken to be 0.1. Rest is similar to UCB implementation.

## Thompson-sampling

Initial number of success and failures for all the arms are taken to be zero. A random number is drawn for each arm from the Beta distribution using numpy as

$$Beta(a, b) = Beta(\#success + 1, \#failures + 1)$$

The arm for which the value drawn is maximum is chosen.

## Thompson-sampling with hint

In this algorithm, the permutation of true means is known. Based on that the maximum value of the true mean can be known. Let two numbers be  $\epsilon_{close}$  to each other if the absolute difference between them is less than epsilon. If the empirical mean of some arm is  $\epsilon_{close}$  to the maximum true mean, that arm gets sampled. Otherwise, the algorithm samples the arms as per thompson-sampling. Epsilon value can be reduced for better accuracy.

## Why it performs better than thompson-sampling?

The algorithm eliminates the cases when an arm with better empirical mean is rejected based on the samples drawn from the  $Beta(a, b)$  distribution.  $Beta(a, b)$  distribution for an arm  $i$  can generate larger value for an arm  $j$ , with very low probability, even if its reward percentage is better than arm  $j$ .

## Generating Rewards

Reward from a chosen arm is based on a fair coin-toss simulated by drawing a number from a uniform distribution (say  $n$ ). Let  $\hat{p}$  be the true mean of the chosen arm. If  $n < \hat{p}$ , a reward of 1 is awarded. This same scheme is followed for all the algorithms.

## Miscellaneous

- To display true means, empirical means, and the number of pulls of an arm, run the file *bandit.py* with the argument *--verbose*. By default, verbosity is turned off.
- The arm with the lower arm number is given priority in case of ties.
- The file to which output is written is labelled as *outputData.txt*. This file is generated in the same directory as *bandit.py*.
- Comments have been provided wherever seemed appropriate

## Plots




Figure 1: Comparison of regret for the four algorithms on instance  $i_1$




Figure 2: Comparison of regret for the four algorithms on instance  $i_2$




Figure 3: Comparison of regret for the four algorithms on instance  $i_3$

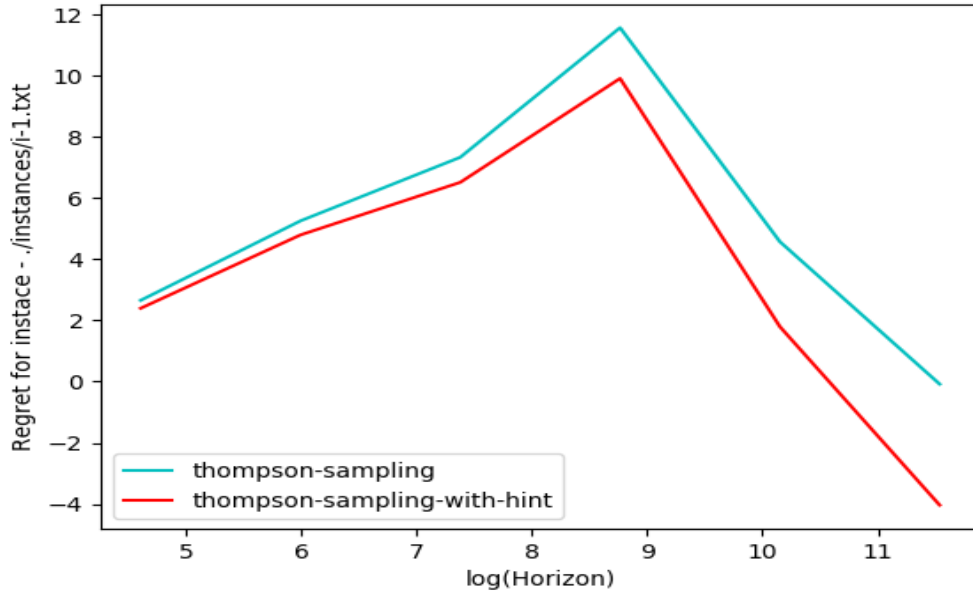


Figure 4: Comparison between Thompson-sampling with and without hint on instance  $i_1$




Figure 5: Comparison between Thompson-sampling with and without hint on instance  $i_2$




Figure 6: Comparison between Thompson-sampling with and without hint on instance  $i_3$

---

## Conclusions

The regret for epsilon-greedy algorithm increases fastest as the horizon increases when compared to other algorithms. Thompson-sampling and KL-UCB, on the other hand, have a much better performance in terms of regret. The regret for UCB, on an average, falls between these two cases. This can be seen with the plots above.

## Explanation of the above behaviour

Owing to the fact that KL-UCB and Thompson-sampling have logarithmic bounds on the regret asymptotically, they are the best performers. Though Thompson-sampling-with-hint outperforms them all. UCB also achieves logarithmic bound, though the bound is not better than KL-UCB. Epsilon greedy is not able to achieve logarithmic regret.

---

## References

- [1] Python documentation available at <https://docs.python.org/3/>
  - [2] Matplotlib documentation available at <https://matplotlib.org/>
  - [3] Numpy documentation available at <https://numpy.org/doc/>
-